

# Clustering & Principal Component Analysis

AERO 689: Machine Learning for Aerospace Engineers

Raktim Bhattacharya

Texas A&M University

# Overview

# Topics Covered

---

1. Classification vs Clustering
2. Clustering Algorithms
3. Cluster Validation
4. Principal Component Analysis (PCA)
5. PCA Applications
6. Integration with Clustering

# Classification vs Clustering

# The Fundamental Difference

---

## Classification (Week 4): Supervised learning

- Requires labeled training data
- Predicts known categories
- Goal: Generalization to new examples
- Example: “Is this bearing fault type A, B, or C?”

## Clustering (Week 5): Unsupervised learning

- No labels required
- Discovers hidden patterns
- Goal: Finding natural groupings
- Example: “What natural fault patterns exist?”

# Mathematical Distinction

---

**Classification:** Learn  $f : X \rightarrow Y$  where  $Y = \{0, 1, \dots, K - 1\}$  is known

- Training data:  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Labels  $y_i$  are given
- Optimize to minimize prediction error

**Clustering:** Find  $f : X \rightarrow \{1, 2, \dots, k\}$  discovered from data

- Training data:  $x_1, x_2, \dots, x_n$  (no labels!)
- Number of clusters  $k$  may be unknown
- Optimize to find natural groups

# When to Use Each Method

---

## Use Classification when:

- You have labeled examples
- Categories are well-defined (taxi, takeoff, cruise)
- Regulatory requirements demand specific fault detection
- Problem: Requires expensive labeling effort

## Use Clustering when:

- No labeled data available
- Exploring new operational regimes
- Discovering anomalies never seen before
- Initial data exploration before labeling

# Application Context

---

## Classification Example (Week 4):

- Trained on known fault types (normal, fault\_a, fault\_b, fault\_c)
- Predicts fault type on new sensor data
- Requires historical fault database

## Clustering Example (Week 5):

- Discovers natural patterns from unlabeled data
- Finds new degradation modes without prior knowledge
- Explores operational envelope
- Creates labels for future classification

**Key Insight:** Clustering creates labels; Classification uses labels

# Decision Framework

---

**Do you have labeled data?**

- **YES** → Classification
  - Logistic regression
  - Decision trees
  - SVM
  - Neural networks

**Do you have labeled data?**

- **NO** → Clustering
  - K-means
  - DBSCAN
  - Hierarchical
  - Gaussian Mixture Models

**Validation Challenge:** How do we know clustering is “correct” without ground truth?

# **Clustering Fundamentals**

# What is Clustering?

**Goal:** Partition data into groups (clusters) where:

- Points in same cluster are similar
- Points in different clusters are dissimilar

**Mathematical Formulation:**

Minimize within-cluster variance:

$$\min_{C_1, \dots, C_k} \sum_{j=1}^k \sum_{x_i \in C_j} \|x_i - \mu_j\|^2$$

where  $\mu_j = \frac{1}{|C_j|} \sum_{x_i \in C_j} x_i$  is cluster center

# K-Means Algorithm

## Lloyd's Algorithm:

1. **Initialize:** Randomly select  $k$  cluster centers  $\mu_1, \dots, \mu_k$
2. **Assignment:** Assign each point to nearest center

$$c_i = \arg \min_j \|x_i - \mu_j\|^2$$

3. **Update:** Recompute centers as mean of assigned points


$$\mu_j = \frac{1}{|C_j|} \sum_{x_i \in C_j} x_i$$

4. **Repeat** steps 2-3 until convergence

**Convergence:** Guaranteed (objective function monotonically decreases)

**Limitation:** May converge to local minimum (solution: run multiple times)

# K-Means Strength: Well-Separated Clusters



**When K-means works perfectly:**

- ✓ Spherical clusters
- ✓ Well-separated
- ✓ Similar sizes
- ✓ Similar densities

**Result:** Perfect match between true clusters and K-means assignments

# K-Means Strength: Code


```
from sklearn.cluster import KMeans
from sklearn.datasets import make_blobs

# Generate well-separated clusters
X, y_true = make_blobs(n_samples=300, centers=4,
                       cluster_std=0.6, random_state=42)

# Apply K-means
kmeans = KMeans(n_clusters=4, random_state=42, n_init=10)
y_pred = kmeans.fit_predict(X)

# Plot comparison
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 4))
ax1.scatter(X[:, 0], X[:, 1], c=y_true, cmap='viridis')
ax2.scatter(X[:, 0], X[:, 1], c=y_pred, cmap='viridis')
ax2.scatter(kmeans.cluster_centers_[:, 0],
            kmeans.cluster_centers_[:, 1],
            marker='X', s=200, c='red', edgecolors='black')
```

# K-Means Weakness 1: Non-Spherical Shapes



**Problem:** K-means assumes spherical clusters

**Failure Mode:**

- ✗ Curved/non-convex shapes
- ✗ K-means draws straight boundaries
- ✗ Incorrectly splits moon shapes

**Solution:** Use DBSCAN or hierarchical clustering

# K-Means Weakness 1: Code


```
from sklearn.datasets import make_moons

# Generate moon-shaped clusters
X_moons, y_moons = make_moons(n_samples=300, noise=0.05,
                                random_state=42)

# K-means fails on non-spherical shapes
kmeans = KMeans(n_clusters=2, random_state=42, n_init=10)
y_pred = kmeans.fit_predict(X_moons)

# Plot
fig, (ax1, ax2) = plt.subplots(1, 2)
ax1.scatter(X_moons[:, 0], X_moons[:, 1], c=y_moons)
ax1.set_title('True Clusters')
ax2.scatter(X_moons[:, 0], X_moons[:, 1], c=y_pred)
ax2.set_title('K-means Result')
```

# K-Means Weakness 2: Different Densities



**Problem:** K-means assumes equal variance

**Failure Mode:**

- Different cluster spreads
- Large variance cluster dominates
- Boundary misplaced

**Solution:**

Standardize features: transform each feature to mean=0, std=1

$$z = \frac{x - \mu}{\sigma}$$

```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

Or use GMM (handles different covariances)


## K-Means Weakness 2: Code

```
# Clusters with different variances
X1 = np.random.normal(0, 1, (200, 2))      # Std = 1
X2 = np.random.normal(5, 3, (200, 2))      # Std = 3
X3 = np.random.normal([0, 6], 0.5, (200, 2)) # Std = 0.5
X = np.vstack([X1, X2, X3])

# K-means struggles with different variances
kmeans = KMeans(n_clusters=3, random_state=42)
y_pred = kmeans.fit_predict(X)

# Large variance cluster dominates
plt.scatter(X[:, 0], X[:, 1], c=y_pred, cmap='viridis')
```

# K-Means Weakness 3: Initialization Sensitivity



**Problem:** Random initialization → different local minima

**Failure Mode:**

- Different results each run
- Some worse than others
- WCSS varies (Within-Cluster Sum of Squares)
  - **Lower WCSS** = better clustering (tighter clusters)
  - **Higher WCSS** = worse clustering (spread out)
- Different seeds → different WCSS values

**Observation:** Seeds 0,2,4 good; Seeds 1,3,5 suboptimal

**Solution:** Use `n_init=10` (run 10 times, keep best)

## K-Means Weakness 3: Code

```
# Demonstrate initialization sensitivity
X, _ = make_blobs(n_samples=300, centers=3, cluster_std=1.0)

# Run with different random seeds
fig, axes = plt.subplots(2, 3, figsize=(12, 8))
for i, ax in enumerate(axes.flat):
    km = KMeans(n_clusters=3, init='random',
                 n_init=1, random_state=i)
    y = km.fit_predict(X)
    ax.scatter(X[:, 0], X[:, 1], c=y)
    ax.scatter(km.cluster_centers_[:, 0],
               km.cluster_centers_[:, 1],
               marker='X', s=200, c='red')
    ax.set_title(f'WCSS={km.inertia_.0f}')

# Solution: n_init=10 (default) runs 10 times, keeps best
```

# **Cluster Validation**

# How Do We Know Clustering is Good?

---

**Challenge:** No ground truth labels!

**Validation Strategies:**

1. **Internal validation:** Use only the data
  - Silhouette score
  - Within-cluster sum of squares (WCSS)
  - Davies-Bouldin index
2. **External validation:** Compare with known labels (if available)
  - Adjusted Rand Index (ARI)
  - Normalized Mutual Information (NMI)
3. **Domain validation:** Domain expertise
  - Do clusters make physical sense?
  - Can we interpret clusters meaningfully?

# Elbow Method

**Goal:** Choose optimal number of clusters  $k$

**Procedure:**


1. Run K-means for  $k = 2, 3, \dots, K_{\max}$
2. Compute Within-Cluster Sum of Squares (WCSS):

$$\text{WCSS}(k) = \sum_{j=1}^k \sum_{x_i \in C_j} \|x_i - \mu_j\|^2$$

3. Plot WCSS vs  $k$
4. Look for “elbow” where improvement diminishes

**Intuition:** Balance fit quality vs model complexity

# Elbow Method: Visualization



## Interpretation:

- **Before  $k=7$ :** Steep decrease (adding clusters helps)
- **At  $k=7$ :** Elbow (optimal)
- **After  $k=7$ :** Flat (diminishing returns)

**Rule:** Choose  $k$  at the elbow

**Application:**  $k=7$  optimal clusters for this dataset

# Elbow Method: Code

```
# Elbow Method implementation
X_scaled = StandardScaler().fit_transform(X)

# Compute WCSS for different k
wcss = []
for k in range(2, 12):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_)

# Plot elbow curve
plt.plot(range(2, 12), wcss, 'bo-')
plt.axvline(x=7, color='red', linestyle='--')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS')
```

**Interpretation:** Elbow at  $k = 7$

# Silhouette Analysis


**Silhouette Score:** How well does point  $i$  fit its cluster?

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))} \in [-1, 1]$$

- $a(i)$ : Mean distance to same-cluster points
- $b(i)$ : Mean distance to nearest other cluster

**Interpretation:**  $s \approx 1$  = well clustered,  $s < 0$  = wrong cluster

# Silhouette Score: Visualization



**Result:**  $s = 0.713$  (good clustering)

**Interpretation:**

- $s > 0.7$ : Strong structure
- $0.5 < s < 0.7$ : Reasonable
- $0.25 < s < 0.5$ : Weak
- $s < 0.25$ : No structure

**Application:** Points with  $s < 0$  are ambiguous cases

# Silhouette Score: Code


```
from sklearn.metrics import silhouette_score

# Compute average silhouette score
s_avg = silhouette_score(X_scaled, clusters)
print(f"Silhouette Score: {s_avg:.3f}")

# Good: s > 0.7
# Reasonable: 0.5 < s < 0.7
# Weak: s < 0.5
```

**Interpretation:** Points with  $s < 0$  are ambiguous conditions

# Silhouette Plot Visualization



## Reading the plot:

- **Width:** Cluster size
- **Beyond avg:** Well-separated
- **Below avg:** Poorly separated
- **Negative:** Wrong cluster

**Observation:** All clusters have  $s > 0.5$  (good)

**Application:** Detect ambiguous cases

# Silhouette Plot: Code

```
from sklearn.metrics import silhouette_samples

# Compute per-sample silhouette scores
s_samples = silhouette_samples(X_scaled, clusters)

# Plot silhouette for each cluster
for k in range(n_clusters):
    cluster_vals = s_samples[clusters == k]
    cluster_vals.sort()
    ax.fill_betweenx(y_range, 0, cluster_vals)

ax.axvline(x=s_avg, color='red', linestyle='--')
```

**Reading:** Wide bars = large clusters,  $s > 0$  = well-clustered

# Validation Against Ground Truth

When labels are available, compare:

Adjusted Rand Index (ARI):

$$\text{ARI} = \frac{\text{RI} - \text{Expected RI}}{\text{Max RI} - \text{Expected RI}}$$

- **Range:**  $[-1, 1]$ , where 1 = perfect match
- **Advantage:** Corrects for chance

```
from sklearn.metrics import adjusted_rand_score

ari = adjusted_rand_score(true_labels, clusters)
print(f"ARI: {ari:.3f}")

# ARI = 1.0: Perfect clustering
# ARI = 0.0: Random clustering
# ARI < 0.0: Worse than random
```

# Case Study: K-Means on Bearing Fault Data

# Case Study: CWRU Bearing Dataset

---

**Dataset:** Case Western Reserve University Bearing Data Center

## Data Details:

- Vibration signals from accelerometers
- Sampling rate: 12 kHz
- Normal + various fault conditions
- MATLAB .mat format

## Our Sample:

- 4 Normal bearings
- 60 Faulty bearings
- 7 extracted features
- **Highly imbalanced!**

**Goal:** Can K-means discover Normal vs Fault without labels?

# True Labels (Supervised View)




Figure 1: True Labels in PCA Space

**Ground Truth** (4 fault types):

- **Blue circles:** Normal (4)
- **Red squares:** Inner Race (12)
- **Green triangles:** Ball (12)
- **Orange diamonds:** Outer Race (36)

**Observation:**

- Normal samples cluster tightly (bottom-left)
- Fault types overlap significantly
- Rightmost points = highest vibration (severe faults)

**Challenge:** Can K-means find these groups?

# K-Means Results ( $k=2$ )




Figure 2: K-Means Clustering

## K-Means Found:

- **Cluster 0** (60 samples): Normal + most faults
- **Cluster 1** (4 samples): Severe faults only

## Metrics:

- Silhouette: 0.66 (good!)
- WCSS: 246.2

**Problem:** K-means grouped by **severity**, not by fault presence!

# Why K-Means Failed

## The Data Imbalance Problem:

Class	Count	Percentage
Normal	4	6.25%
Fault	60	93.75%

## What K-Means saw:


- 4 extreme outliers (severe faults with very high RMS)
- 60 “similar” samples (normal + mild/moderate faults)

## What K-Means did:

- Cluster 0: Everything with lower vibration
- Cluster 1: Only the 4 severe faults


**Lesson:** K-means finds density patterns, not our desired labels!

# Comparison: True vs K-Means



**True Labels**

4 fault types (colors = fault location)



**K-Means (k=2)**

2 clusters (by vibration severity)

**Key Insight:** Clustering finds patterns that **exist in feature space**, not necessarily the patterns we want!

# Lessons Learned

---

## When K-Means Works:

- Balanced cluster sizes
- Well-separated clusters
- Spherical cluster shapes

## When K-Means Struggles:

- **Imbalanced data** (like our 4 vs 60)
- Feature overlap between classes
- Non-spherical cluster shapes

## Better Approaches for Fault Detection:

- **Anomaly detection**: Train on normal data, flag outliers
- **Supervised learning**: Use labels when available
- **More data**: Collect balanced training set

# K-Means Limitations

---

## Assumptions:

- Spherical clusters (assumes equal variance)
- Same cluster sizes
- Well-separated clusters

## Sensitive to:

- Outliers (pull cluster centers)
- Initialization (local minima)
- Feature scaling (always standardize!)

## When K-means fails:

- Non-spherical shapes
- Varying densities
- Highly imbalanced data

# Other Clustering Methods

Algorithm	Key Idea	When to Use
<b>DBSCAN</b>	Density-based regions	Unknown $k$ , outlier detection
<b>Hierarchical</b>	Build cluster tree	Explore hierarchy, dendrogram
<b>GMM</b>	Probabilistic Gaussians	Soft cluster assignments

## Quick Summary:

- **DBSCAN**: Groups dense regions, labels outliers as noise (-1)
- **Hierarchical**: Agglomerative (bottom-up) or Divisive (top-down)
- **GMM**: Mixture of Gaussians with soft probabilities

**For this course:** Focus on K-means; other methods follow similar principles



# **Principal Component Analysis**

# Dimensionality Reduction Motivation

---

## The Curse of Dimensionality:

- Modern sensor systems: 100+ features
- High-dimensional data is hard to visualize
- Computational cost grows exponentially
- Many features are correlated (redundant)

**Goal:** Reduce dimensions while preserving information

**PCA:** Find directions of maximum variance

**Example:** 6 feature dataset → 3 principal components

# PCA: The Big Picture

---

**Problem:** Data in  $\mathbb{R}^d$  (e.g.,  $d = 6$  sensors)

**Goal:** Project to  $\mathbb{R}^k$  where  $k < d$  (e.g.,  $k = 3$ )

**Constraint:** Minimize information loss

**PCA Solution:** Project onto directions of maximum variance

**Key Insight:** Variance = information content

- High variance dimensions contain signal
- Low variance dimensions contain noise

# PCA Mathematics

**Given:** Data matrix  $X \in \mathbb{R}^{n \times d}$  ( $n$  samples,  $d$  features)

**Step 1:** Center the data

$$X' = X - \bar{X} \quad \text{where } \bar{X} = \frac{1}{n} \sum_{i=1}^n x_i$$

**Step 2:** Compute covariance matrix

$$\Sigma = \frac{1}{n} X'^T X' \in \mathbb{R}^{d \times d}$$

**Step 3:** Eigendecomposition

$$\Sigma v_j = \lambda_j v_j$$

where  $v_j$  = eigenvector (principal component),  $\lambda_j$  = eigenvalue

# PCA: Principal Components

---

**Eigenvectors**  $v_1, \dots, v_d$ :

- Directions of maximum variance
- Orthogonal to each other
- Ordered:  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$

**Variance Explained** by PC  $j$ :  $\frac{\lambda_j}{\sum_i \lambda_i}$

**Projection**:  $Z = X'V_k$  where  $V_k = [v_1, \dots, v_k]$

# Geometric Interpretation

---

**PCA finds coordinate system aligned with data variance**


- **PC1:** Direction of maximum spread
- **PC2:** Direction of maximum spread orthogonal to PC1
- **PC3:** Direction of maximum spread orthogonal to PC1 and PC2
- ...

**Projection:** Rotate axes to align with principal directions

**Dimensionality Reduction:** Keep only top  $k$  axes

**Information Loss:** Variance in discarded components

# Scree Plot Visualization



## Interpretation:

- **PC1-3:** Capture most variance
- **PC4-6:** Diminishing returns
- **PC7-10:** Noise

**Rule:** Keep until 95% cumulative variance

**Result:** Keep 4 components (10D → 4D)

**Application:** 3-4 components typically capture system dynamics

# Scree Plot: Code

```
from sklearn.decomposition import PCA

# Fit PCA
pca = PCA().fit(X_scaled)
var_ratio = pca.explained_variance_ratio_

# Plot scree
plt.bar(range(1, len(var_ratio)+1), var_ratio)
plt.xlabel('Principal Component')
plt.ylabel('Variance Ratio')

# Cumulative variance
cumsum = np.cumsum(var_ratio)
n_components = np.argmax(cumsum >= 0.95) + 1
print(f"{n_components} components for 95% variance")
```

**Rule:** Keep components until 95% cumulative variance

# Cumulative Variance

**Choose number of components:** Find where cumulative variance reaches threshold

```
cumsum = np.cumsum(pca.explained_variance_ratio_)
n_components_95 = np.argmax(cumsum >= 0.95) + 1
print(f"{n_components_95} components explain 95% of variance")
```

**Typical Result:** 3-4 components for 95% variance

**Application:** Often 3-4 components capture 95% of system dynamics

# Connection to Spectral Analysis

## Eigenanalysis Review:

### Spectral Decomposition:

$$A\phi_j = \lambda_j\phi_j$$

- $\phi_j$ : Eigenvector
- $\lambda_j$ : Eigenvalue

**Insight:** PCA finds “data modes” via eigenanalysis

### PCA:

$$\Sigma v_j = \lambda_j v_j$$

- $v_j$ : Principal component
- $\lambda_j$ : Variance

# Spectral Decomposition Analogy

Generic Spectral Analysis	PCA
Eigenvector $\phi_j$	Principal component $v_j$
Eigenvalue $\lambda_j$	Variance $\lambda_j$
Spectral contribution	Variance explained
Reduced-order model	Dimensionality reduction
Neglect small eigenvalues	Neglect low-variance components

**Both:** Find directions that capture system behavior efficiently

**Insight:** PCA finds “data modes” via eigenanalysis

# PCA Applications

# Feature Reduction

---

**Problem:** System has 6 correlated features

- feature\_1
- feature\_2
- feature\_3
- feature\_4
- feature\_5
- feature\_6

**Question:** Can we reduce to 3 features without losing information?

**Solution:** Apply PCA, keep top 3 components

# PCA Implementation

```
# Full feature set
features = ['feature_1', 'feature_2', 'feature_3', 'feature_4',
            'feature_5', 'feature_6']
X = data[features].values

# Standardize
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# PCA to 3 components
pca = PCA(n_components=3)
X_pca = pca.fit_transform(X_scaled)

print(f"Original dimensions: {X_scaled.shape[1]}")
print(f"Reduced dimensions: {X_pca.shape[1]}")
print(f"Variance preserved: {pca.explained_variance_ratio_.sum():.1%}")
```

**Result:** 6D → 3D, preserving 95%+ variance

# 2D PCA Visualization

Project to 2D and visualize with cluster colors:

```
pca_2d = PCA(n_components=2)
X_pca_2d = pca_2d.fit_transform(X_scaled)

plt.figure(figsize=(10, 8))
for label in data['label'].unique():
    mask = data['label'] == label
    plt.scatter(X_pca_2d[mask, 0], X_pca_2d[mask, 1],
                label=label, alpha=0.6, s=30)
plt.xlabel(f'PC1 ({pca_2d.explained_variance_ratio_[0]:.1%})')
plt.ylabel(f'PC2 ({pca_2d.explained_variance_ratio_[1]:.1%})')
plt.legend()
```

**Observation:** Classes naturally separate in PC space!

# Feature Interpretation: Loadings

**Loadings:** Weights of original features on components

```
loadings = pd.DataFrame(  
    pca_2d.components_.T,  
    columns=['PC1', 'PC2'],  
    index=features  
)  
print(loadings.round(3))
```

**Example Interpretation:**

Feature	PC1	PC2	Interpretation
feature_1	0.52	-0.18	PC1 = primary axis
feature_2	0.48	-0.22	PC1 = primary axis
feature_3	0.31	0.58	PC2 = secondary axis
feature_5	0.27	0.64	PC2 = secondary axis

# Preserving Classification Performance

---

**Critical Question:** Does PCA preserve classification capability?

**Approach:**

1. Train classifier on **full features**
2. Train classifier on **PCA-reduced features**
3. Compare accuracy

# Preserving Classification Performance (contd.)

```
from sklearn.linear_model import LogisticRegression

# Full feature classifier
clf_full = LogisticRegression()
clf_full.fit(X_scaled, fault_labels)
acc_full = clf_full.score(X_test_scaled, fault_test)

# PCA-reduced classifier
X_pca = pca.fit_transform(X_scaled)
clf_pca = LogisticRegression()
clf_pca.fit(X_pca, fault_labels)
acc_pca = clf_pca.score(X_test_pca, fault_test)

print(f"Full: {acc_full:.3f} | PCA: {acc_pca:.3f}")
```

**Goal:** Minimal accuracy loss with major dimension reduction

# Information Loss Analysis

**Reconstruction Error:** Measure information lost in dimensionality reduction

```
# Reconstruct from reduced representation
X_reconstructed = pca.inverse_transform(X_pca)

# Overall error
recon_error = np.mean((X_scaled - X_reconstructed)**2)
print(f"MSE: {recon_error:.4f}")

# Per-feature error
for i, feature in enumerate(features):
    feature_error = np.mean((X_scaled[:, i] -
                            X_reconstructed[:, i])**2)
    print(f"{feature}: {feature_error:.4f}")
```

**Decision:**

- Accept error if **critical features** preserved
- Monitor reconstruction error for important features
- Higher tolerance for redundant features

# PCA for Noise Filtering

**Insight:** Noise projects onto low-variance components

**Denoising Procedure:**

1. Apply PCA to noisy data
2. Keep only high-variance components
3. Reconstruct from reduced representation

**Signal is in high-variance components, noise in low-variance**

```
# Denoise by keeping top k components
pca_denoise = PCA(n_components=3)
X_denoised = pca_denoise.inverse_transform(
    pca_denoise.fit_transform(X_noisy)
)
```

**Application:** Clean noisy sensor data

## **Integration: PCA + Clustering**

# Why Combine PCA and Clustering?

---

## Motivation:

- High-dimensional data is computationally expensive
- Correlated features confuse clustering algorithms
- Visualization impossible in  $>3$  dimensions

## Solution Pipeline:

1. Apply PCA:  $d$  dimensions  $\rightarrow k$  dimensions
2. Apply clustering on reduced space
3. Interpret results in original space

**Benefits:** Faster, more stable, visualizable

# Implementation Pipeline


```
# Step 1: PCA
pca = PCA(n_components=3)
X_pca = pca.fit_transform(X_scaled)

# Step 2: Cluster
kmeans = KMeans(n_clusters=7, random_state=42)
clusters = kmeans.fit_predict(X_pca)

# Step 3: Validate
from sklearn.metrics import silhouette_score
print(f"Silhouette: {silhouette_score(X_pca, clusters):.3f}")
```

**Result:** 6D → 3D, then K-means clustering

# Visualization in PC Space



## Pipeline:

1. Start: 6D data
2. PCA: Reduce to 2D
3. K-means: Cluster in 2D
4. Visualize: See all clusters

**Advantage:** Visualize high-dimensional clustering!

**Application:** View classes in 2D PC space

# PCA + Clustering: Code

```
# Pipeline: PCA then clustering
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

# Cluster in reduced space
kmeans = KMeans(n_clusters=5, random_state=42)
clusters = kmeans.fit_predict(X_pca)

# Visualize
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=clusters)
plt.xlabel(f'PC1 ({pca.explained_variance_ratio_[0]:.1%})')
plt.ylabel(f'PC2 ({pca.explained_variance_ratio_[1]:.1%})')
```

**Advantage:** Visualize high-dimensional clustering in 2D!

# Parameter Sensitivity

## How do choices affect results?

```
for n_comp in [2, 3, 4, 5]:  
    X_pca = PCA(n_components=n_comp).fit_transform(X_scaled)  
    clusters = KMeans(n_clusters=7).fit_predict(X_pca)  
    ari = adjusted_rand_score(true_labels, clusters)  
    print(f"Components: {n_comp}, ARI: {ari:.3f}")
```

**Observation:** Performance plateaus after 3-4 components

# Clustering Algorithm Comparison

```
algorithms = {
    'K-means': KMeans(n_clusters=7),
    'DBSCAN': DBSCAN(eps=0.5, min_samples=10),
    'GMM': GaussianMixture(n_components=7)
}
for name, model in algorithms.items():
    clusters = model.fit_predict(X_pca)
    print(f"{name}: ARI={adjusted_rand_score(y, clusters):.3f}")
```

**Result:** K-means often best for structured data classification

# Real-World Case Study

---

**Scenario:** Large dataset with 10,000 samples, 12 features

**Goal:** Identify operational patterns for optimization

**Approach:**

1. **Data collection:**  $10,000 \text{ samples} \times 12 \text{ features} \times 7200 \text{ timesteps} = 864\text{M data points}$
2. **PCA:** Reduce 12D  $\rightarrow$  4D (preserving 96% variance)
3. **Clustering:** K-means with  $k = 5$  operational modes
4. **Analysis:** Characterize each mode by context variables
5. **Optimization:** Recommend best practices from efficient clusters

**Result:** 2-3% efficiency improvement = significant cost savings

# **Supervised vs Unsupervised**

# The Complete ML Taxonomy

Method	Type	Input	Output	Example
<b>Regression</b>	Supervised	$X, y$ continuous	$\hat{y} \in \mathbb{R}$	Predict drag coefficient
<b>Classification</b>	Supervised	$X, y$ categorical	$\hat{y} \in \{0, 1, \dots, K - 1\}$	Identify fault type
<b>Clustering</b>	Unsupervised	$X$ only	Cluster assignments	Discover patterns
<b>PCA</b>	Unsupervised	$X$ only	Reduced $X'$	Reduce features

# Decision Flowchart

---

**Labeled Data?** → **Supervised**

- Continuous output → Regression
- Categorical output → Classification

**Hybrid:** Use clustering to create labels, then train classifier

**No Labels?** → **Unsupervised**

- Find groups → Clustering
- Reduce dimensions → PCA

# When to Use Each

---

**Regression:** Continuous predictions

**Classification:** Categorical decisions

**Key Question:** Do you have labels?

- Yes → Supervised (Regression/Classification)
- No → Unsupervised (Clustering/PCA)

**Clustering:** Discover patterns

**PCA:** Reduce dimensions

# Wrap-Up

# Key Takeaways

---

## Classification vs Clustering:

- Classification uses labels; Clustering discovers labels
- Both partition data, different paradigms

## Clustering Algorithms:

- K-means: Fast, requires  $k$ , spherical clusters
- DBSCAN: Finds outliers, arbitrary shapes, needs  $\varepsilon$
- Hierarchical: Explores all  $k$ , dendrogram
- GMM: Probabilistic, soft assignments

## Validation:

- Elbow method, silhouette score (internal)
- ARI, NMI (external, when labels available)
- Always check domain interpretability

# Key Takeaways (continued)

---

## PCA:

- Projects to directions of maximum variance
- Connection to modal analysis (eigenvectors!)
- Choose  $k$  to preserve 95% variance

## Integration:

- PCA before clustering: faster, more stable
- Safety-aware reduction: preserve classification performance
- Visualization: cluster in 3D PC space

## Application Context:

- Unsupervised learning discovers unknown failure modes
- PCA reduces feature redundancy
- Always validate physical interpretability

# Homework 5 Preview

**Assignment:** Data Clustering & PCA Analysis

## Part 1: Clustering (35 points):

- Implement K-means with elbow method selection
- Apply DBSCAN for anomaly detection
- Hierarchical clustering with dendrogram
- Validate against ground truth labels

## Part 2: PCA (35 points):

- Full PCA with scree plot analysis
- Reduce 6D → 3D feature space
- Classification-aware validation (preserve performance)
- Interpret principal components physically

# Homework 5 Preview (continued)

---

## Part 3: Integration (30 points):

- PCA + Clustering pipeline
- Compare performance: full vs reduced space
- Visualize 3D clustering results
- Technical writeup with domain interpretation

## Bonus (10 points):

- Gaussian Mixture Model implementation
- Advanced visualization (interactive 3D plots)
- Novel application

**Due:** Next week

**Dataset:** Synthetic multi-sensor data (provided)

# Next Week: Week 6

---

**Coming Up:** Advanced Dimensionality Reduction

**Topics:**

- Proper Orthogonal Decomposition (POD)
- Connection between PCA and POD
- t-SNE for visualization
- Autoencoders (neural network approach)
- Applications to CFD data compression

**Bridge to Deep Learning:** Autoencoders lead to Week 7 (Neural Networks)

# Questions & Discussion

---

## Open Floor:

- Clustering vs classification confusion?
- PCA mathematics clarification?
- Applications?
- Homework questions?

**Office Hours:** See syllabus

## Resources:

- Bishop PRML Chapter 9 (Mixture Models)
- Bishop PRML Chapter 12 (PCA)
- Class notes on Canvas

# Thank You!

---

## AERO 689: Machine Learning for Aerospace Engineers

Dr. Raktim Bhattacharya  
Texas A&M University

**Next Lecture:** Week 6 - Advanced Dimensionality Reduction